

# Deploying Django web applications in Docker containers

**Jamie Hewland**  
PyConZA 2017  
PRAEKELT.ORG



00

Hi 🖐️, I'm Jamie

**SRE at Praekelt.org**



**@JayH5**





**@jayhewland**



**/jhewland**

**We moved to containers because we needed to deploy a lot of sites quickly**


**...and manually managing Python processes across many servers wasn't scaling.**



springster


[Home](#)
[My Future](#)
[My Body](#)

[Girl Stories](#)
[My Life](#)

NEW LOOK...  
STILL FABULOUS.  
CHECK US OUT.






**Tick Tock. Tick Tock. Why is my biological...**

Read this piece below to...

[#puberty](#)




**The problem with dandruff**

Keep a level head

[#relationshine](#)

[English](#) | [français](#)




[Home](#) | [Pregnancy](#) | [Baby](#) |
   
[Staying healthy](#) |
   
[About BabyCenter](#)

Recommended

[Staying healthy](#)

**Warning signs:** When to get medical help



[Pregnancy](#)

**What is Zika?:** Zika is an infection from mosquitoes

Pregnancy


LANGUE: **FRANÇAIS**

[CHANGER LA LANGUE](#) ▼

INTERNET OF GOOD THINGS

[Inscrivez-vous](#)
[Recherche](#)

[Menu](#)



**DROITS DE L'ENFANT**

**Personne n'a le droit de me faire**

# PRAEKELT ORG

# Deploying Django web applications in *Docker containers*

Specifically, Docker containers to run under a *container orchestration* system.

- What are Docker containers?\*
- What is container orchestration?\*
- How is Django typically deployed?
- How do I deploy Django in a Docker container?

\*Very roughly speaking

# 01

## Containers & container orchestration

It's 

# What's a (Linux) container?

Isolation of a process' view of its operating environment (via *namespaces*)

- Process trees
- User IDs
- Networking
- Mounted file systems...

Limitation/prioritization of resources (via *cgroups*)

- CPU
- Memory
- Block I/O
- Networking...

# Docker containers



Docker is the most popular container technology.

- Batteries included
- Easy-to-use, lots of sensible defaults
- “*Copy-on-write union filesystem*”: containers can start up very quickly and share a lot of image data
- Images available for all the software you know & love



```
> $ docker run ubuntu  
debian  
redis  
rabbitmq  
python  
postgres  
sentry  
nginx  
nodejs
```

# Docker terminology

file

layers on disk

running process

**Docker-  
file**

`docker build`

**Docker  
image**

`docker run`

**Docker  
container**

`ls Dockerfile`  
`ls *.dockerfile`

`docker images`

`docker ps`

**Terms “*image*” and “*container*” often conflated**

# Why containers?

## Consistent portability

- A clean way to package software
- With (almost) everything it needs to run
- With a single, simple entry-point
- Limit access to resources
- Eliminates “but it works on *my* machine”

# Container orchestration

“Container Orchestration refers to the automated arrangement, coordination, and management of software containers.”

*Container Orchestration with Kubernetes: An Overview* - Adrian Chifor  
<http://bit.ly/2takgmd>

# Container orchestration

Achieved through a variety of services:

- Service discovery
- Load balancing
- Health checks
- Resource management
- More...

# Container orchestration is 🔥🔥🔥



**kubernetes**

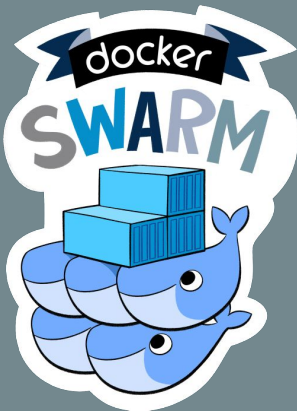


HashiCorp

**Nomad**



**DC/OS**



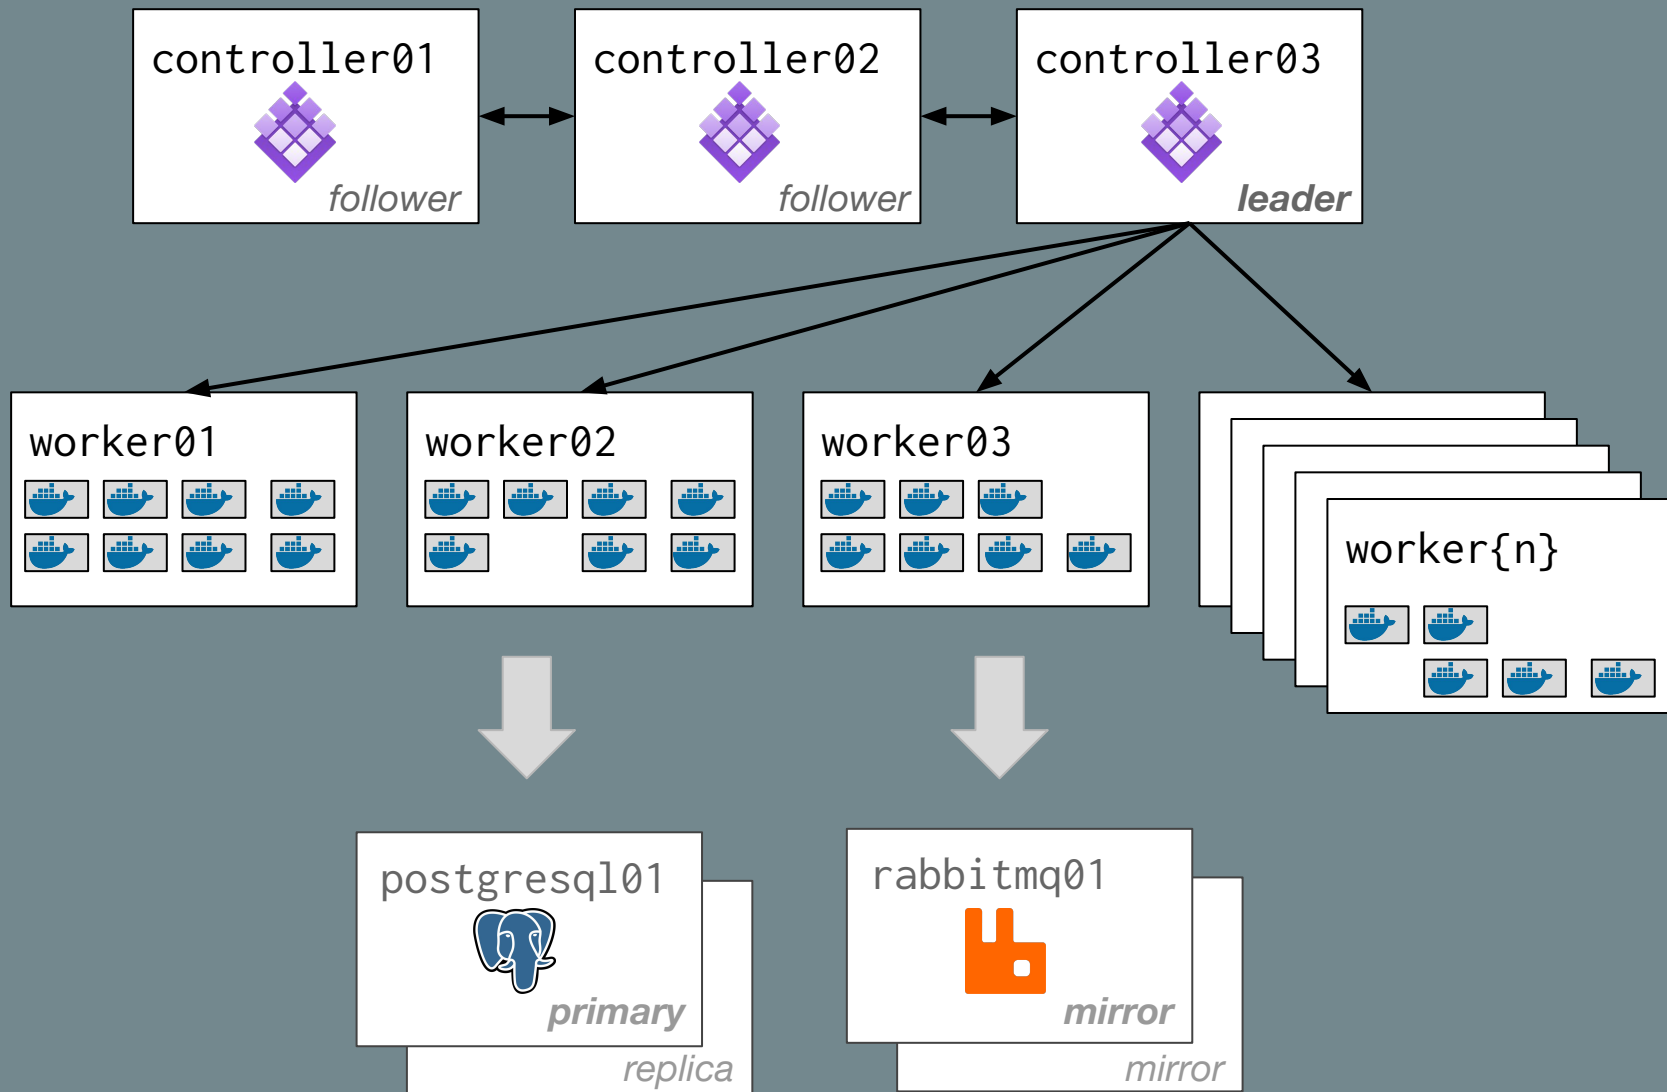
**amazon**  
web services



**Microsoft Azure**



**Google Cloud Platform**



controller01



*follower*

controller02



*follower*

controller03



*leader*

**Control-plane**

worker01



worker02



worker03



worker{n}



**Pool of workers**

**Stateful services**

postgresql01



*primary*

*replica*

rabbitmq01



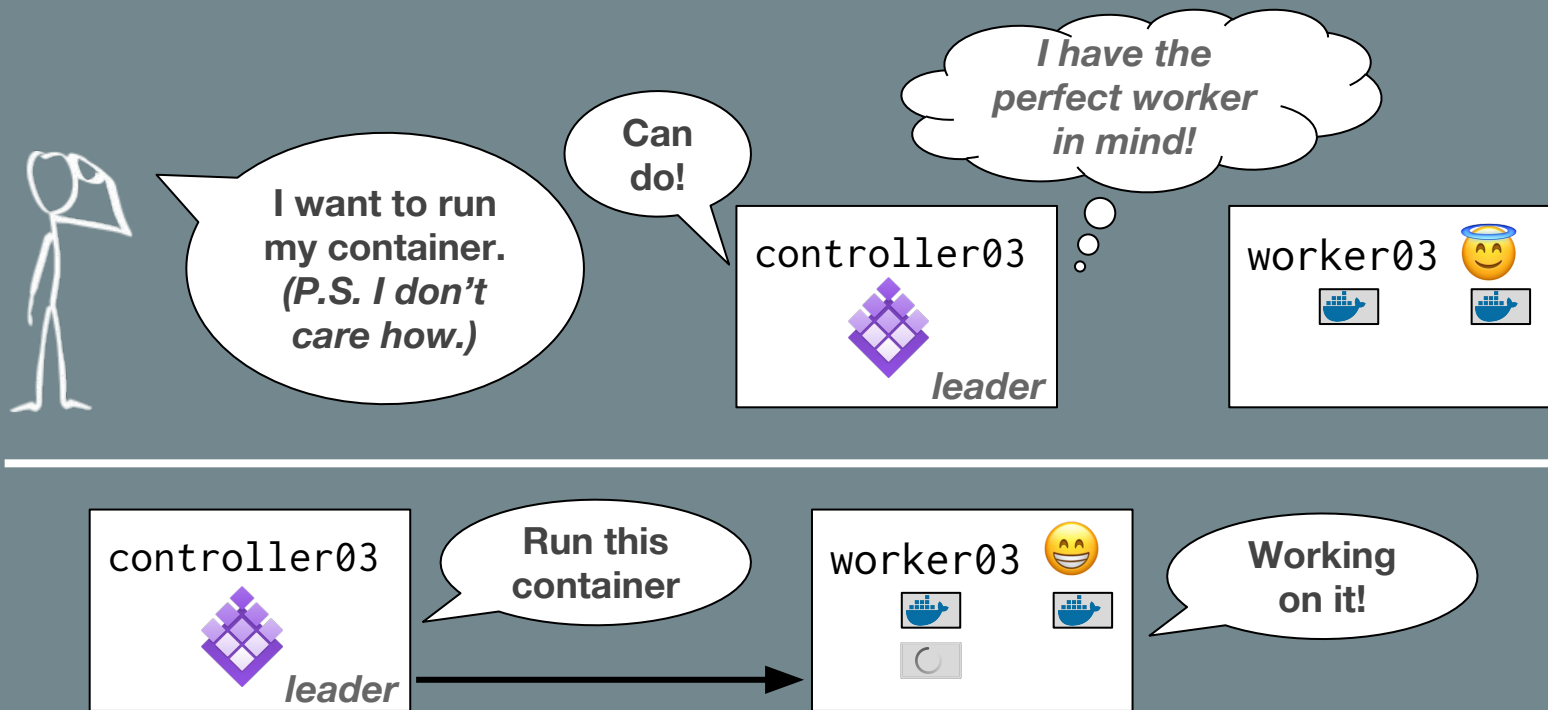
*mirror*

*mirror*



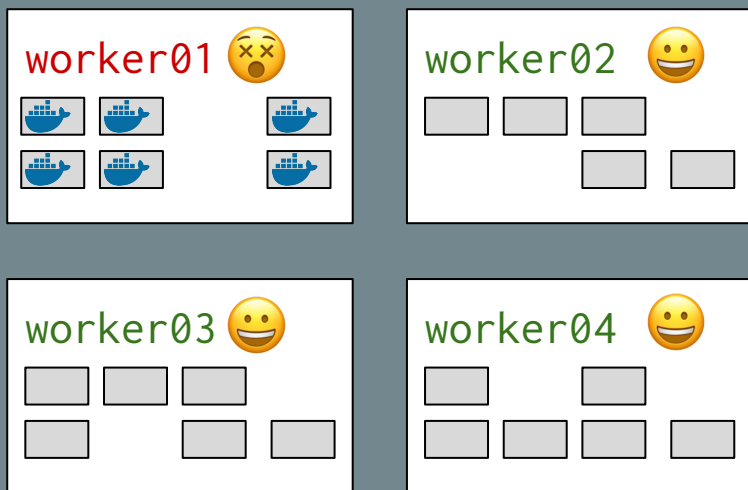
# Advantage 1: Deployments

Don't need to choose how/where to run apps

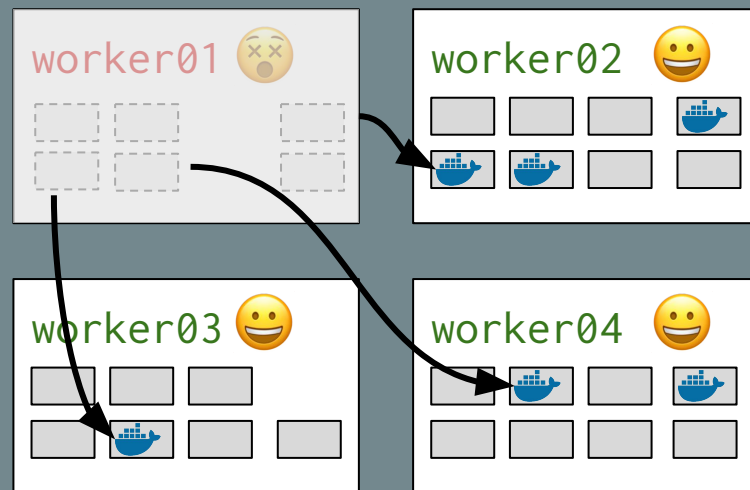


# Advantage 2: Failover

Containers get moved off unhealthy worker hosts



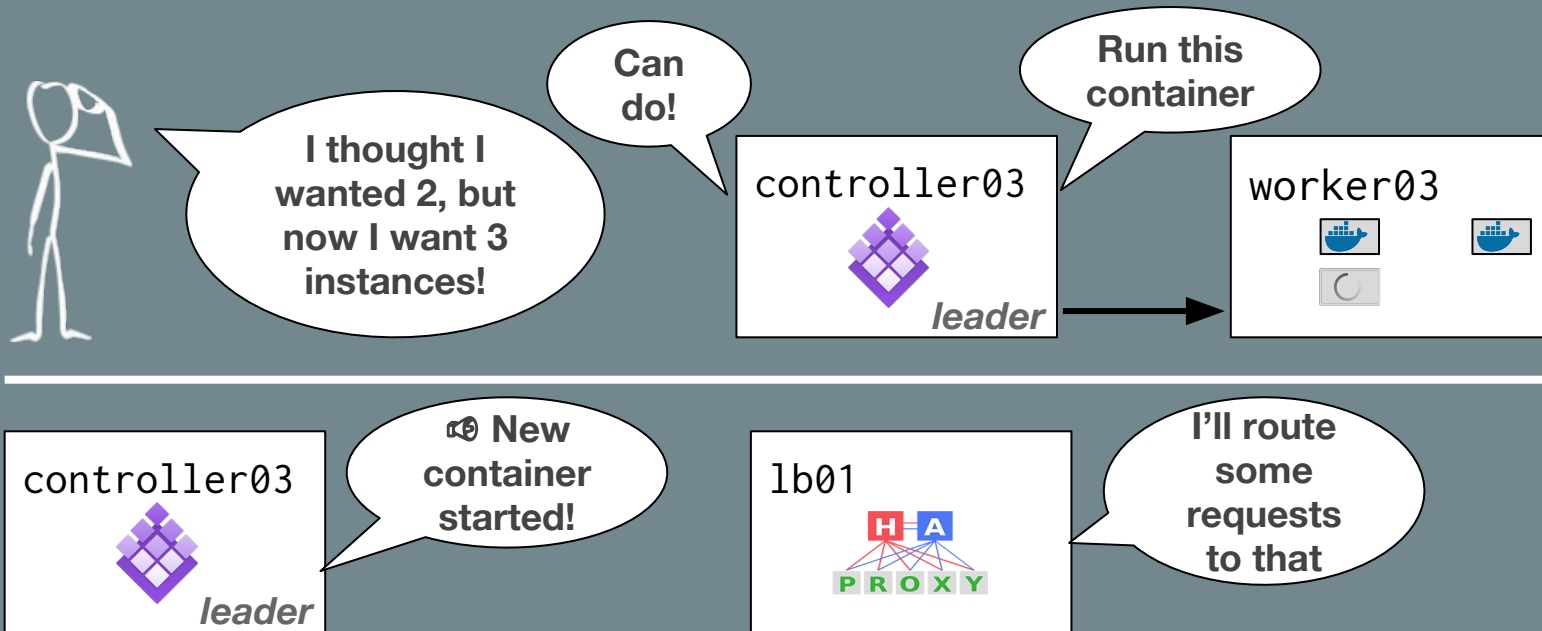
1. worker01 is unhealthy



2. Containers migrated off worker01

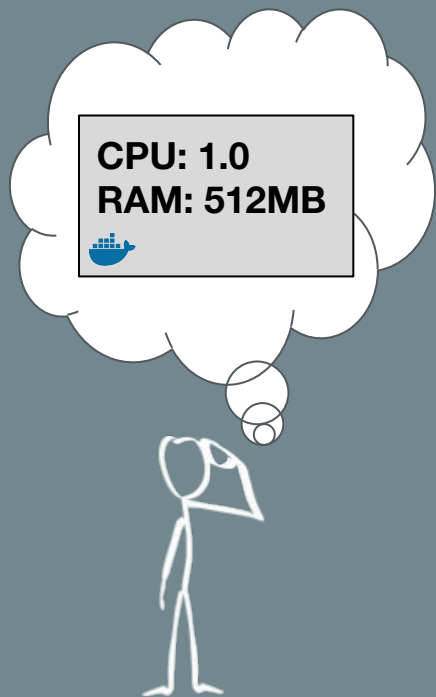
# Advantage 3: Scaling

Can scale the number of running containers



# Advantage 4: Resource utilisation

Orchestration systems can pack containers efficiently



worker03

CPU: 0.4  
RAM: 512MB



CPU: 1.0  
RAM: 768MB



CPU: 0.5  
RAM: 128MB



CPU: 1.9/3.0  
RAM: 1.375/2.0GB

worker05

CPU: 1.0  
RAM: 768MB



CPU: 1.0  
RAM: 512MB



CPU: 0.5  
RAM: 256MB



CPU: 2.5/3.0  
RAM: 1.5/2.0GB

02

# Deploying Django web applications

*“Webservers”*

# django

- Open source web framework
- Very popular
- Roughly model-view-controller (MVC)
- Featureful: Caching, i18n, middleware...
- Lots of 3rd party extensions

# Running Django

- Django applications typically served using a WSGI server
- Web Server Gateway Interface (WSGI): PEP 3333
- Two sides: *“server”/“gateway”* and *“application”/“framework”*
- Server calls application once per request with data from request

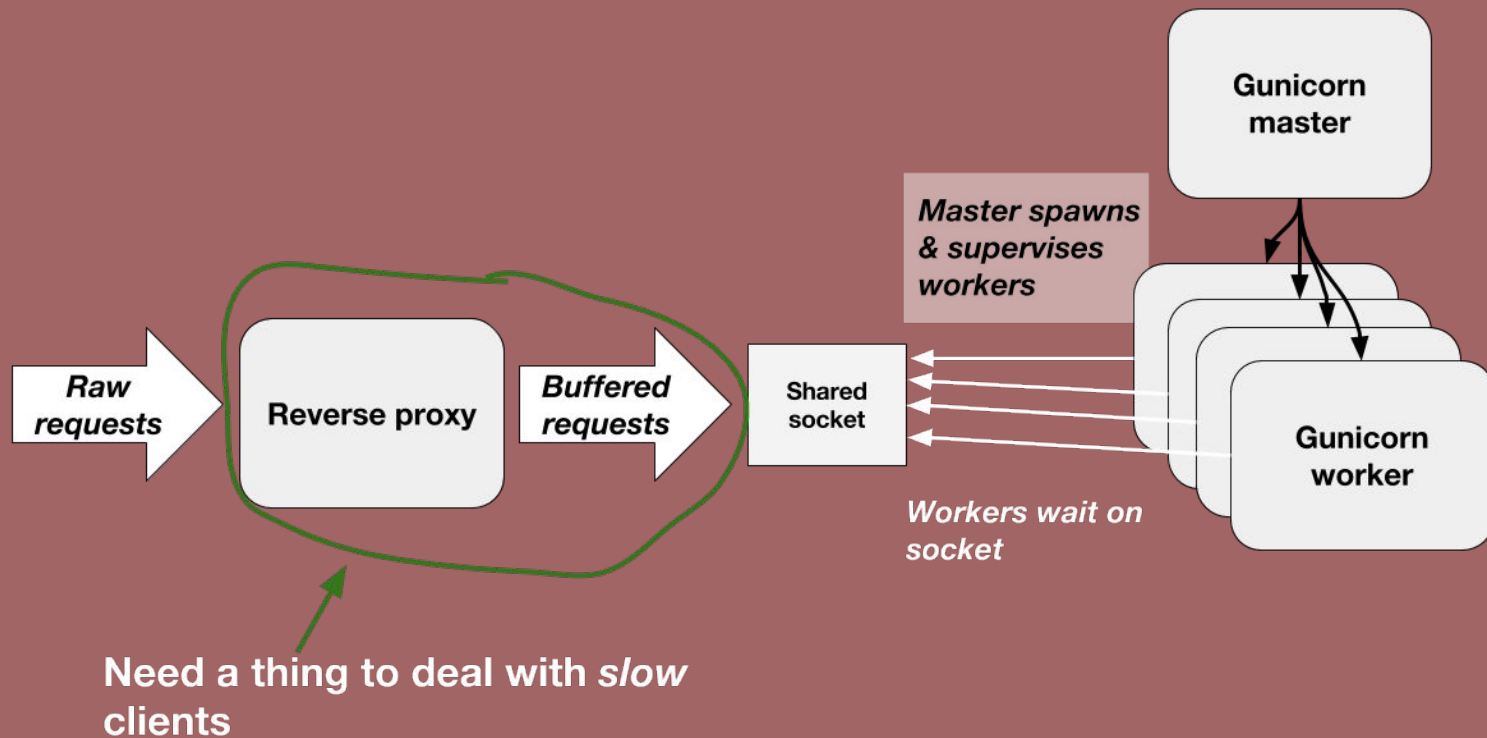


- **Gunicorn is a WSGI server**
- **Based on *Unicorn* (Ruby software)**
- **Pre-fork worker model:**
  - One master process spawns one or more worker processes
  - Master terminates workers if they take too long
  - Workers = 2-4 x number of cores
- **Designed to only serve fast clients\***

*\*With its default worker implementation*



# How it fits together



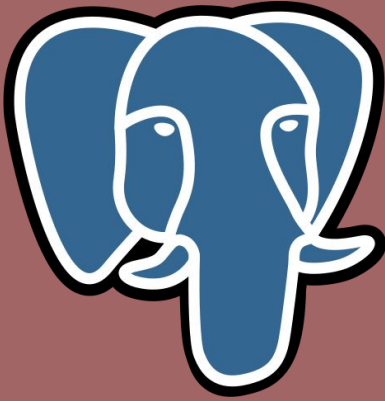


- Reverse proxy most often used with Gunicorn is *Nginx*
- Very fast and battle-tested: “shields” our Python code from the outside
- Can use it to do other useful stuff:
  - Serve static files (CSS, JS, fonts...)
  - Caching
  - *Compression, SSL, more...*



- Django under Gunicorn runs only in response to requests
- What about long-running and/or periodic tasks?
- Celery: Distributed Task Queue
- Integrates with Django
- *But now we need a message broker*

# Other things needed for Django



A database:  
*probably* PostgreSQL

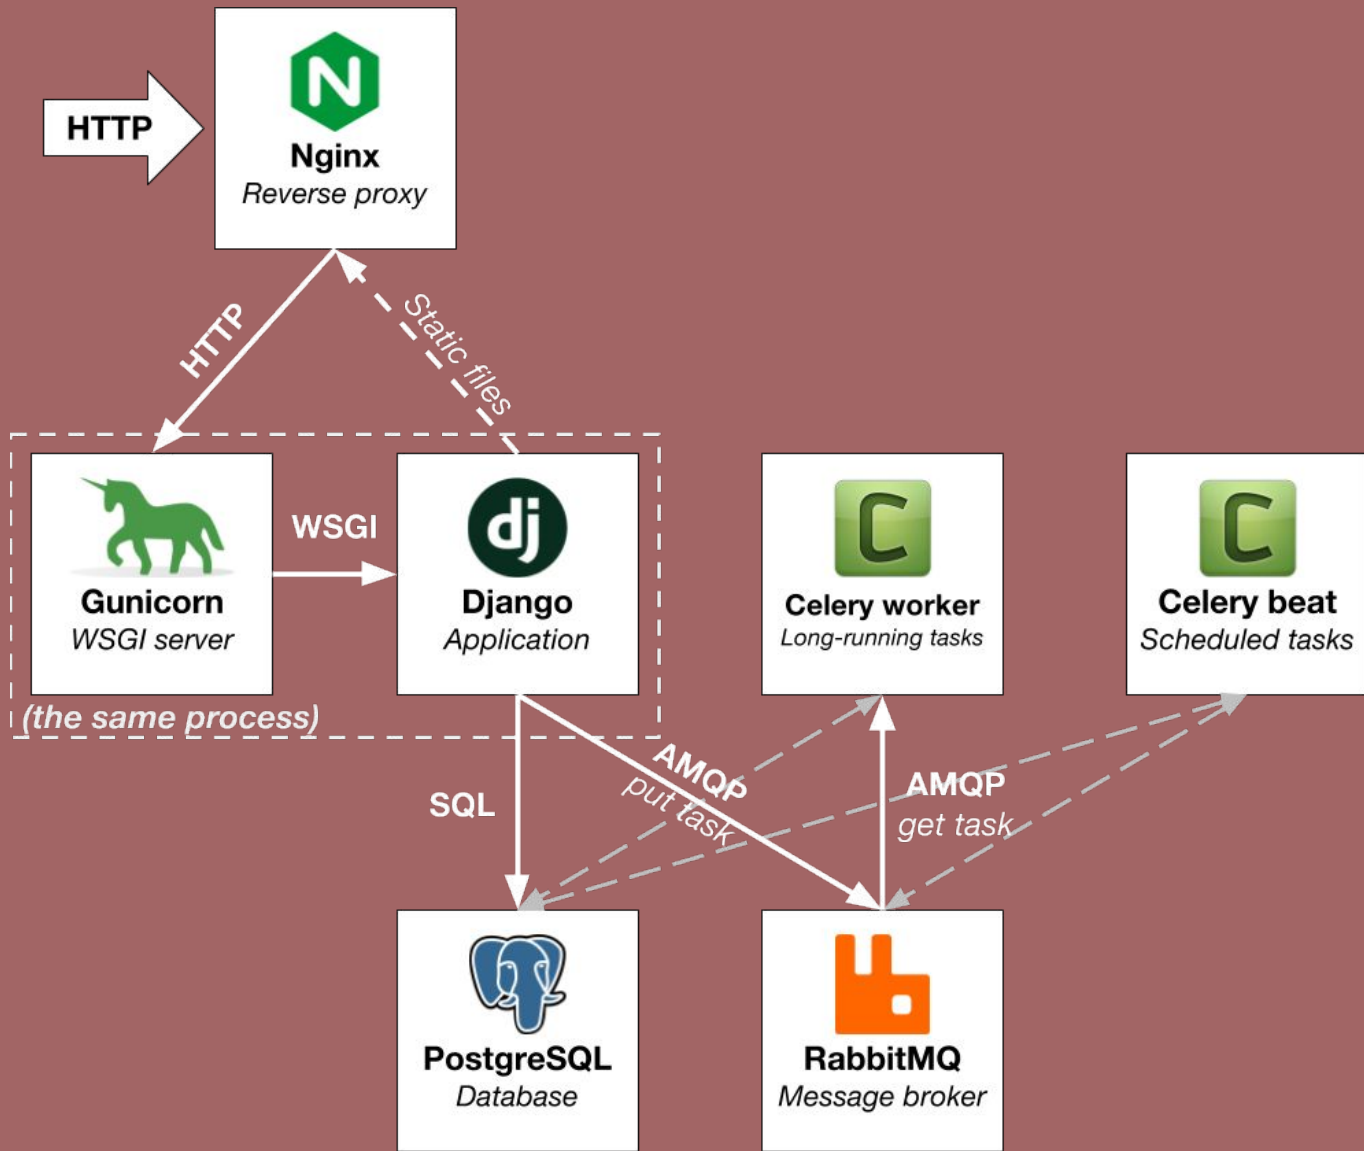


redis

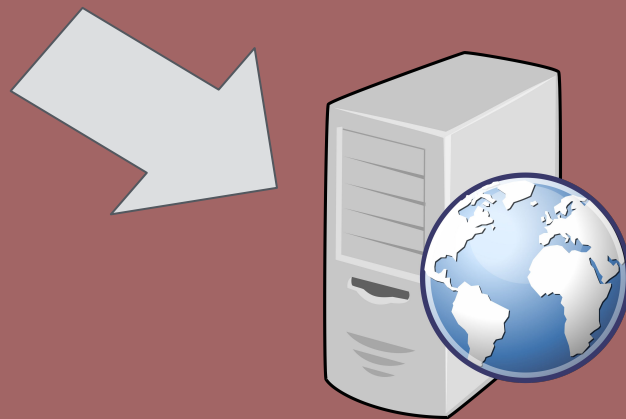
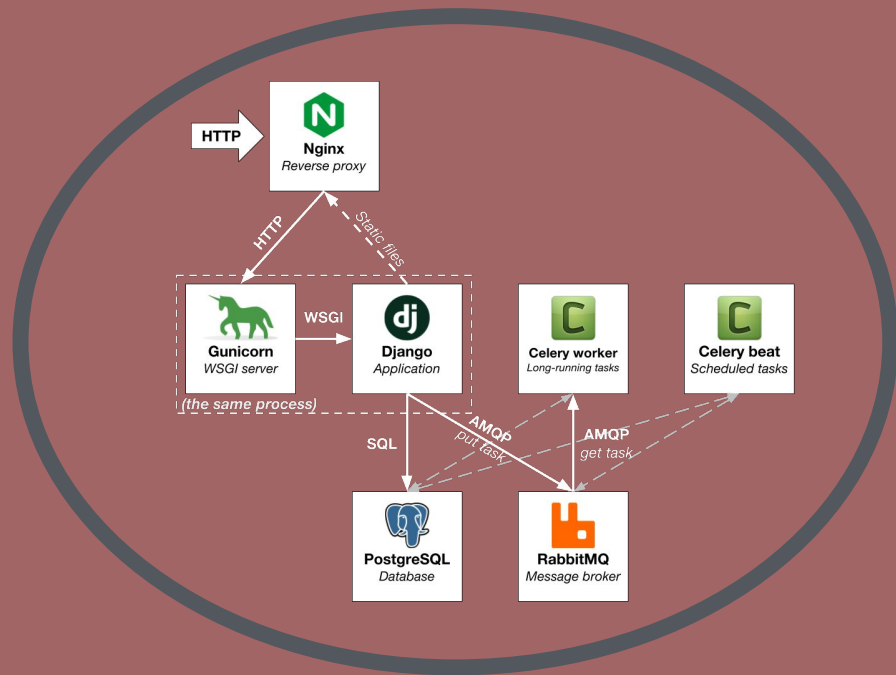


MEMCACHED

(Optional) caching:  
Redis or Memcached



# *Take all of that and...*

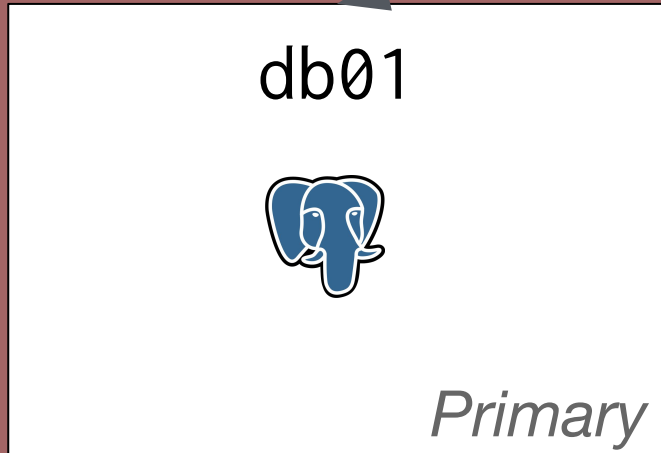


# *...just chuck it on a server*

webserver01



*All the things...*





webserver01



*Some of the things...*

celery01



db01



*Primary*

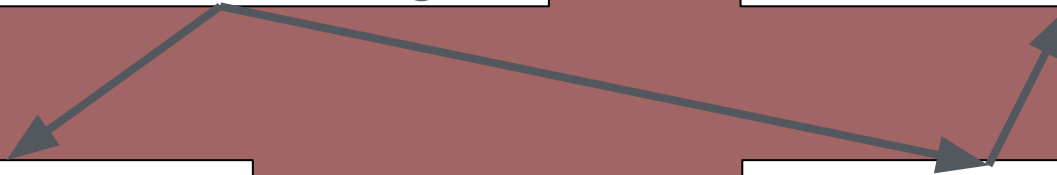
*Replica*

amqp01

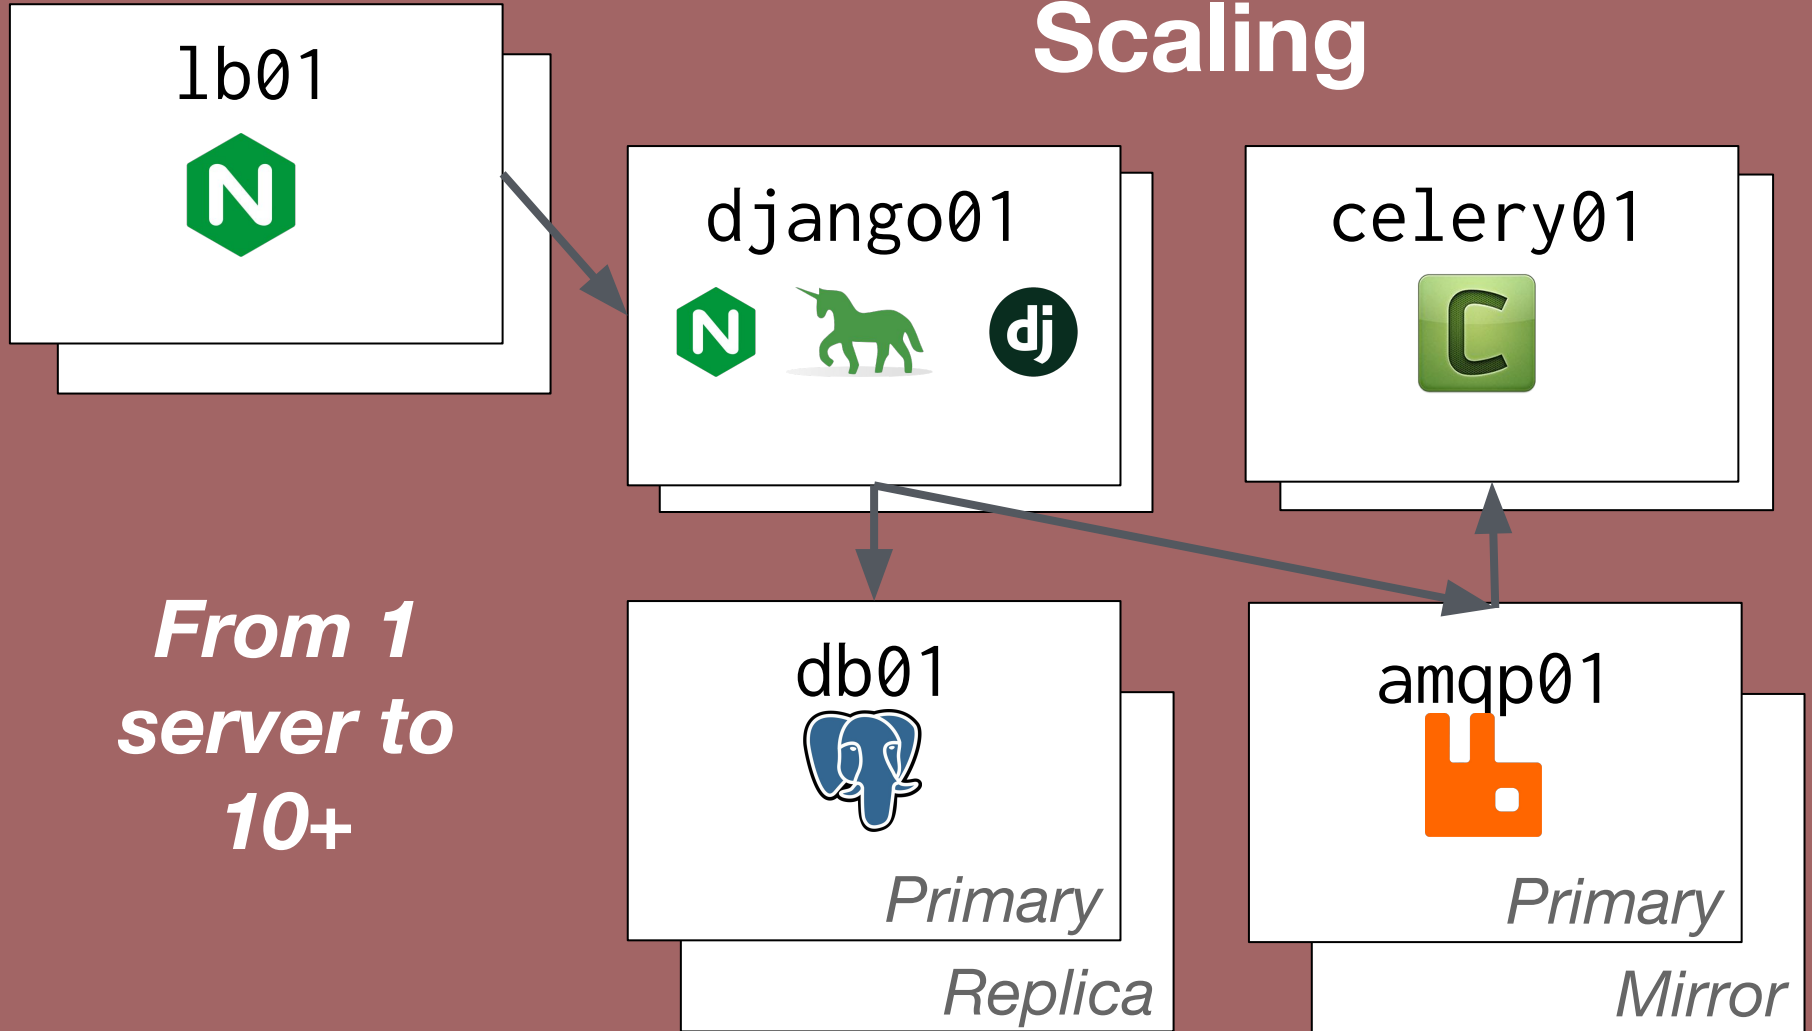


*Primary*

*Mirror*



# Scaling



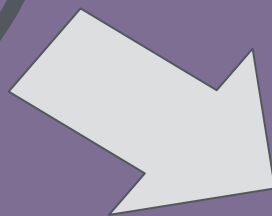
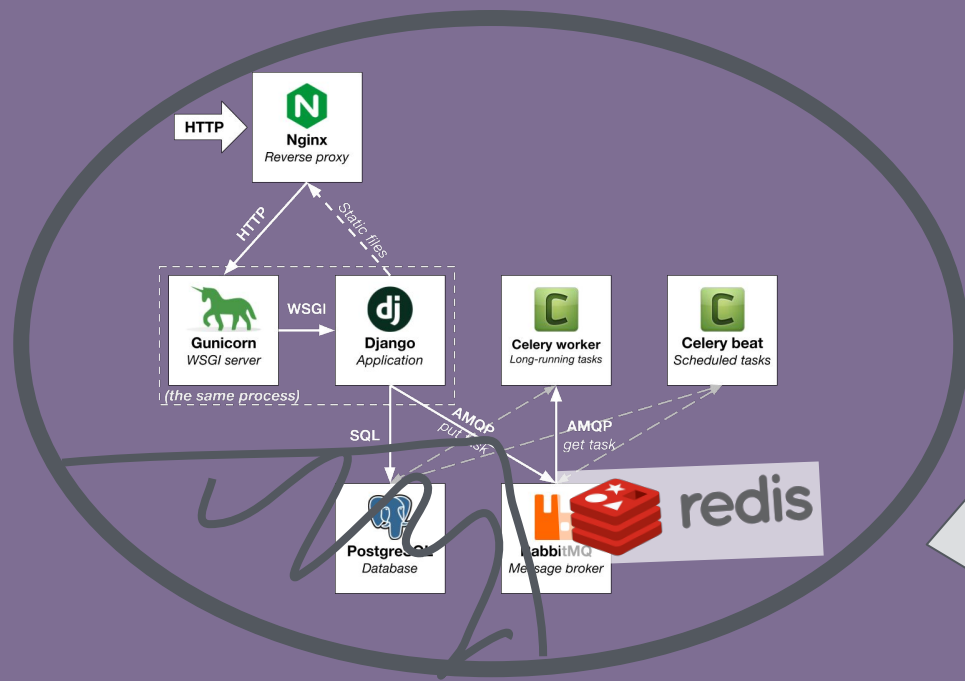
03

# Containerizing Django

**Making it fit**

**You'll never guess what  
we tried first...**

# Take most of that and...

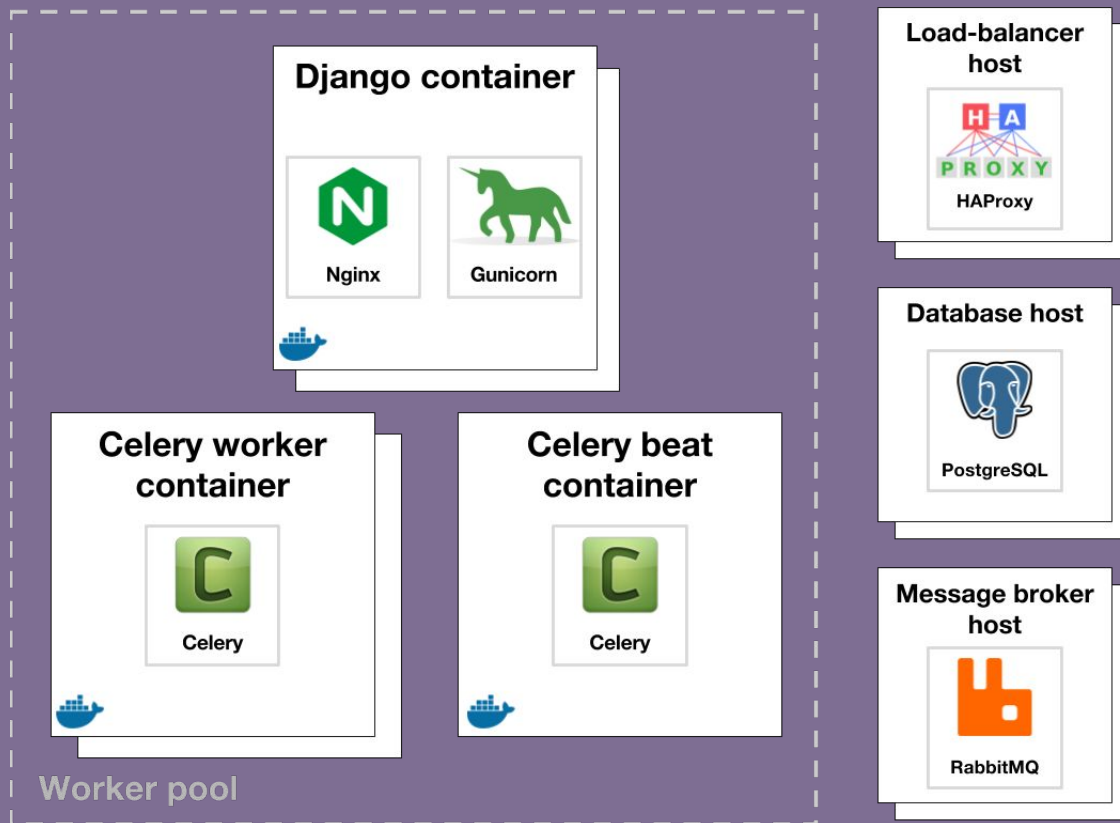


# ...just chunked it in a container

# Do not do this

- Docker containers are *not* mini-VMs
- Isolate processes, not services
- Good health checks very difficult
- No programs to manage multiple processes (no init system)
- Container orchestration systems expect containers to be ephemeral

# What we've settled on



# Configuration

- ***Don't* want to build a Docker image with all the config files inside it**
  - Not flexible, slow reconfiguration
  - Secrets in Docker image
- **Difficult to move configuration files around with containers**
- **Solution: Django settings module reads config from environment variables**



# django-environ

**DATABASE\_URL=postgres://user:pass@db01/dbname**

**CACHE\_URL=memcache://mem01:11211,mem02:11211**

**EMAIL\_URL=smtp+tls://user:pass@smtp01:465**

```
1  import environ
2
3  env = environ.Env()
4
5  # Raises ImproperlyConfigured exception if SECRET_KEY not in os.environ
6  SECRET_KEY = env("SECRET_KEY")
7  DEBUG = env("DEBUG", default=False)
8
9  DATABASES = {"default": env.db(default="sqlite:///tmp/db.sqlite3")}
10 CACHES = {"default": env.cache(default="locmemcache://")}
11 EMAIL_CONFIG = env.email_url(default="consolemail://")
--
```

# Startup (entrypoint) scripts

When the container starts we need to do some things...

- Run database migrations
- Create a superuser account on first run
- Set some default Gunicorn arguments
- Switch to a non-root user
- More... (but hopefully not)

# Logging

Since containers are ephemeral, so are their log files

- Log everything to stdout/stderr
- Container orchestrators will collect this
- Even more important that only one thing runs in a container
- Bonus points: make your logs machine-readable

# User-uploaded files

User-uploaded files in Django can be stored in a “media” directory

- Don't do this (containers or not)
- Extra hard if you have to manage networked storage
- Use `django-storages`, store in S3

# django-bootstrap image

(Not to be confused with CSS Bootstrap)

- Standardized base image for all our Django deployments
- Nginx configuration optimised for Django in a container
- Startup scripts for Django & Celery
- Thoroughly tested with example app



[praekeltfoundation/docker-django-bootstrap](https://github.com/praekeltfoundation/docker-django-bootstrap)

# django-bootstrap Dockerfile

```
1  FROM praekeltfoundation/django-bootstrap
2
3  # Copy in our source code and install it
4  COPY . .
5  RUN pip install -e .
6
7  # Tell Django & Celery where our config & app is
8  ENV DJANGO_SETTINGS_MODULE cake_service.settings
9  ENV CELERY_APP cake_service
10
11 # Collect static files so that they can be served
12 RUN django-admin collectstatic --noinput
13
14 # Pass arguments to Gunicorn: where our WSGI app is
15 CMD ["cake_service.wsgi:application"]
```

04

# Where to from here?

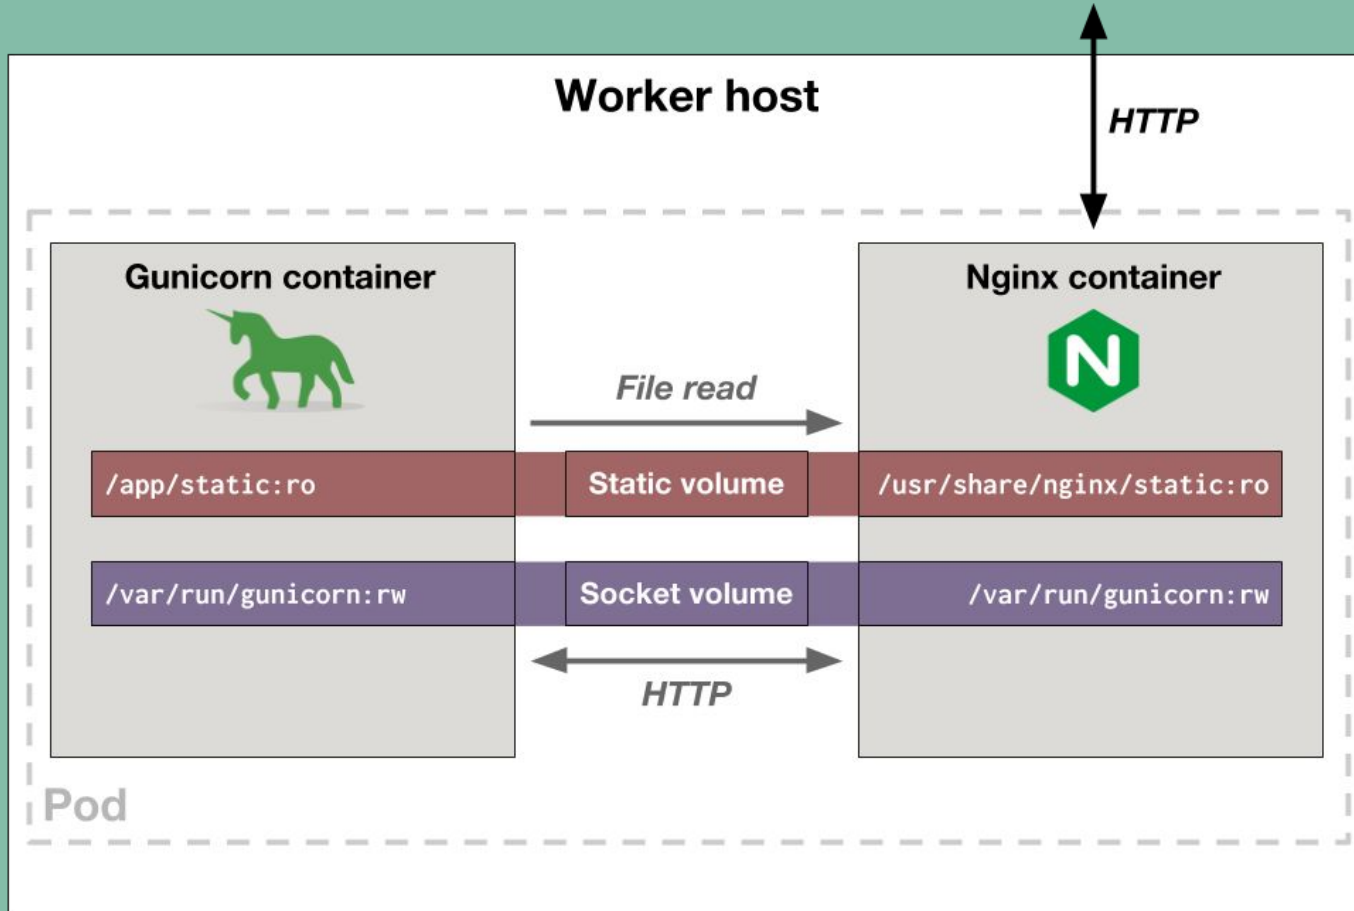
Further improvements

But you're still running more than one thing in a container?





# Deploying as a *pod*



# Configuration secrets

- We're now putting passwords in environment variables
- But env vars are quite easy to leak
- Container orchestration platforms have tools for storing secret data securely
- Dynamic credential management (Hashicorp Vault): credentials only valid as long as the container exists

# Metrics via nginx-lua-prometheus

```
1  # HELP nginx_http_connections Number of HTTP connections
2  # TYPE nginx_http_connections gauge
3  nginx_http_connections{state="reading"} 0
4  nginx_http_connections{state="waiting"} 0
5  nginx_http_connections{state="writing"} 1
6
7  # HELP nginx_http_request_duration_seconds HTTP request latency
8  # TYPE nginx_http_request_duration_seconds histogram
9  nginx_http_request_duration_seconds_bucket{host="localhost",le="00.005"} 3
10 nginx_http_request_duration_seconds_bucket{host="localhost",le="00.010"} 3
11 ...
12 nginx_http_request_duration_seconds_bucket{host="localhost",le="10.000"} 3
13 nginx_http_request_duration_seconds_bucket{host="localhost",le="+Inf"} 3
14 nginx_http_request_duration_seconds_count{host="localhost"} 3
15 nginx_http_request_duration_seconds_sum{host="localhost"} 0.002000093460083
16
17 # HELP nginx_http_requests_total Number of HTTP requests
18 # TYPE nginx_http_requests_total counter
19 nginx_http_requests_total{host="localhost",status="200"} 1
20 nginx_http_requests_total{host="localhost",status="404"} 1
```

# Metrics via nginx-lua-prometheus

- Prometheus can poll container orchestrator to know where to scrape
- Get Django-specific metrics from Nginx (e.g. all requests not to `/static/`)
- “Free” metrics for all apps with same base image
- But you *should* properly instrument your Django application...

05

## In conclusion

Containers are cool, but...

# Containers/container orchestration can seem complicated...

- Persistent storage difficult\*
- No config files\*
- No log files\*
- Can only run one thing per container\*
- Can't SSH in\*
- Distributed system 🥲

*\*Roughly speaking*

**...but they have some big advantages**

- **Easy deployments**
- **Easy scaling**
- **Efficient resource usage**
- **Generally increased automation**
- **Consistently packaged apps**

# Containers for Django

A common base image can provide...

- Tested and optimised server (Nginx) config
- Encapsulate best practices for containers & container orchestration platform
- Consistent platform for deploying apps
- Potential for adding new features



# Questions?



[praeeltfoundation/docker-django-bootstrap](https://github.com/praeeltfoundation/docker-django-bootstrap)

PRAEELT.ORG

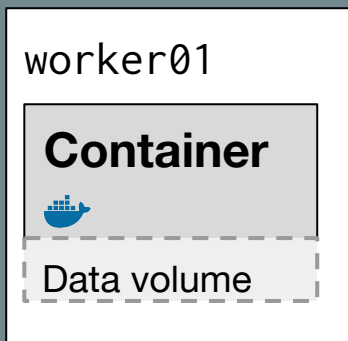
# Thank you.

Special thanks to Jeremy Thurgood (@jerith) for reviewing my code.  
Go see his talk later!

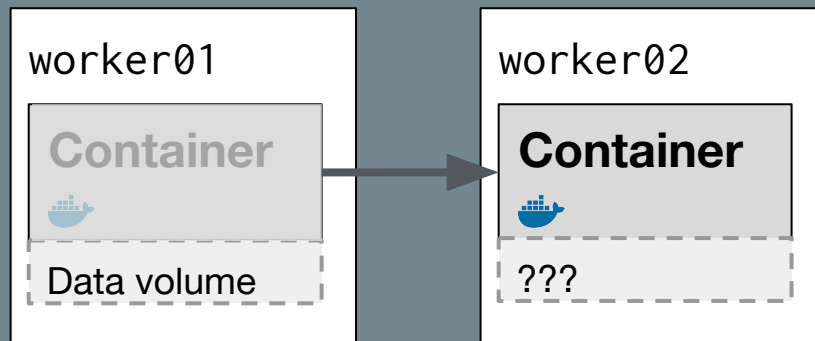
PRAEKELT.ORG

# Complication 1: Persistent storage

Moving data is harder than moving a container



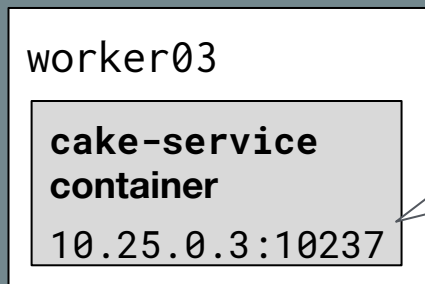
1. The container needs to be moved



2. But its data needs to move with it

# Complication 2: Networking

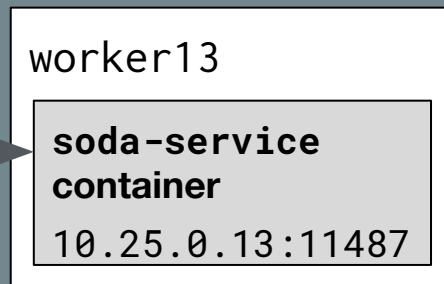
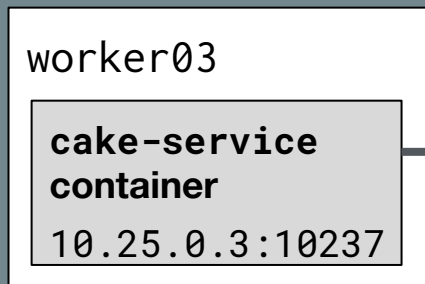
Things move around and thus have weird addresses



I need to speak  
to  
soda-service



Use  
soda-service  
.marathon.141b  
.thisdcos  
.directory



# Complication 3: Debugging

It's hard to just "SSH into" a container

```
curl controller01:8080/v2/apps ...
```



1. Find which worker the container is on

```
docker ps | grep cake-service
```



3. Find the container ID

```
ssh -t public01 ssh worker42
```



worker42



2. SSH into the worker

```
docker exec -it 981681d291ab bash
```

```
root@981681d291ab:~#
```

4. Run Bash in the container